

# Redis 集群操作

Name : 曲中岭

Email:[zlingqu@126.com](mailto:zlingqu@126.com)

Q Q :441869115

## 一、创建集群

### 1.1 规划

10.1.5.205:7001 master --> 10.1.5.205:7004 slave

10.1.5.205:7002 master --> 10.1.5.205:7005 slave

10.1.5.205:7003 master --> 10.1.5.205:7006 slave

### 1.2 创建集群

这里不介绍集群的部署，相关文档有很多，只介绍运维人员可能需要了解的相关集群操作。

```
create      host1:port1 ... hostN:portN  
  
            --replicas <arg>
```

举例：

--replicas 1 每个 master 创建副本的数量是 1

```
redis-trib.rb create --replicas 1 10.1.5.205:7001 10.1.5.205:7002 10.1.5.205:7003
```

10.1.5.205:7004 10.1.5.205:7005 10.1.5.205:7006

这条命令，会默认将前三个设置为 master，后三个设置为 slave，且一一对应，如果没有问题，输入 yes 即可创建完成。

```
[root@hadoop05 redis-cluster]#  
[root@hadoop05 redis-cluster]#  
[root@hadoop05 redis-cluster]# redis-trib.rb create --replicas 1 10.1.5.205:7001 10.1.5.205:7002 10.1.5.205:7003 10.1.5.205:7004 10.1.5.205:7005 10.1.5.205:7006  
>>> Creating cluster  
>>> Performing hash slots allocation on 6 nodes...  
Using 3 masters:  
10.1.5.205:7001  
10.1.5.205:7002  
10.1.5.205:7003  
Adding replica 10.1.5.205:7004 to 10.1.5.205:7001  
Adding replica 10.1.5.205:7005 to 10.1.5.205:7002  
Adding replica 10.1.5.205:7006 to 10.1.5.205:7003  
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001  
slots:0-5460 (5461 slots) master  
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002  
slots:5461-10922 (5462 slots) master  
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003  
slots:10923-16383 (5461 slots) master  
S: bc3e08e793544c02c665fe7200053f86f45bc9bd 10.1.5.205:7004  
replicates da150342ca272529d5dab62a557da5206305c333  
S: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005  
replicates a84292141c294ad66fa0434c4e04c740769285e6  
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006  
replicates faa7e2b785c07ed4b074b485cd78072e76337372  
Can I set the above configuration? (type 'yes' to accept): yes  
>>> Nodes configuration updated  
>>> Assign a different config epoch to each node  
>>> Sending CLUSTER MEET messages to join the cluster  
Waiting for the cluster to join...  
>>> Performing Cluster Check (using node 10.1.5.205:7001)  
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001  
slots:0-5460 (5461 slots) master  
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002  
slots:5461-10922 (5462 slots) master  
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003  
slots:10923-16383 (5461 slots) master  
M: bc3e08e793544c02c665fe7200053f86f45bc9bd 10.1.5.205:7004  
slots: (0 slots) master  
replicates da150342ca272529d5dab62a557da5206305c333  
M: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005  
slots: (0 slots) master  
replicates a84292141c294ad66fa0434c4e04c740769285e6  
M: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006  
slots: (0 slots) master  
replicates faa7e2b785c07ed4b074b485cd78072e76337372
```

## 二、集群信息查看

### 2.1 使用 check 检查集群

当出现下面一行时，证明集群工作正常，可使用 zabbix 等监控此信息来判断集群状态：

[OK] All 16384 slots covered.

同时如果出现：

0 additional replica(s)

证明有从节点挂了或则本身没有配置从节点，此时虽然不影响正常工作，但应该引起管理员注意。

```
[root@hadoop05 redis-cluster]# redis-trib.rb check 10.1.5.205:7001
>>> Performing cluster check (using node 10.1.5.205:7001)
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001
  slots:1366-5461 (4096 slots) master
  1 additional replica(s)
M: 9987adaef51933673afe6f5d7c0ea292c94f2182 10.1.5.205:7007
  slots:0-1365,5462-6826,10923-12287 (4096 slots) master
  0 additional replica(s)
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003
  slots:12288-16383 (4096 slots) master
  1 additional replica(s)
S: bc3e08e793544c02c665fe7200053f86f45bc9bd 10.1.5.205:7004
  slots: (0 slots) slave
  replicates da150342ca272529d5dab62a557da5206305c333
S: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005
  slots: (0 slots) slave
  replicates a84292141c294ad66fa0434c4e04c740769285e6
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002
  slots:6827-10922 (4096 slots) master
  1 additional replica(s)
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006
  slots: (0 slots) slave
  replicates faa7e2b785c07ed4b074b485cd78072e76337372
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
[root@hadoop05 redis-cluster]#
```

当其中一个 master 挂掉，且没有从顶替时，使用 check 查看：

[ERR] Not all 16384 slots are covered by nodes.

```
[root@hadoop05 redis-cluster]# redis-trib.rb check 10.1.5.205:7001
>>> Performing cluster check (using node 10.1.5.205:7001)
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001
  slots:1366-5461 (4096 slots) master
  1 additional replica(s)
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003
  slots:12288-16383 (4096 slots) master
  1 additional replica(s)
S: bc3e08e793544c02c665fe7200053f86f45bc9bd 10.1.5.205:7004
  slots: (0 slots) slave
  replicates da150342ca272529d5dab62a557da5206305c333
S: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005
  slots: (0 slots) slave
  replicates a84292141c294ad66fa0434c4e04c740769285e6
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002
  slots:6827-10922 (4096 slots) master
  1 additional replica(s)
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006
  slots: (0 slots) slave
  replicates faa7e2b785c07ed4b074b485cd78072e76337372
[OK] All nodes agree about slots configuration.
>>> check for open slots...
>>> check slots coverage...
[ERR] Not all 16384 slots are covered by nodes.
```

## 2.2 使用 info 查看集群信息

```
[root@hadoop05 redis-cluster]# redis-trib.rb info 10.1.5.205:7001
10.1.5.205:7001 (da150342...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7007 (9987adae...) -> 0 keys | 4096 slots | 0 slaves.
10.1.5.205:7003 (faa7e2b7...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7002 (a8429214...) -> 0 keys | 4096 slots | 1 slaves.
[OK] 0 keys in 4 masters.
0.00 keys per slot on average.
[root@hadoop05 redis-cluster]#
```

当其中一个 master 挂掉，且没有从顶替时，使用 info 查看：

[ERR] Sorry, can't connect to node 10.1.5.205:7007

```
[root@hadoop05 redis-cluster]# redis-trib.rb info 10.1.5.205:7001
[ERR] Sorry, can't connect to node 10.1.5.205:7007
10.1.5.205:7001 (da150342...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7003 (faa7e2b7...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7002 (a8429214...) -> 0 keys | 4096 slots | 1 slaves.
[OK] 0 keys in 3 masters.
0.00 keys per slot on average.
[root@hadoop05 redis-cluster]#
```

## 2.3 查看 key 的 slot

先使用 CLUSTER KEYSLOT key 计算出 slot，再使用 info 查看对应的 master

```


>>> Check slots coverage...
[OK] All 16384 slots covered.
[root@hadoop05 7004redis]# redis-cli -p 7002 CLUSTER KEYSLOT name
(integer) 5798
[root@hadoop05 7004redis]#
[root@hadoop05 7004redis]#
[root@hadoop05 7004redis]# redis-trib.rb check 10.1.5.205:7002
>>> Performing cluster check (using node 10.1.5.205:7002)
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002
  slots:6827-10922 (4096 slots) master
  1 additional replica(s)
S: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005
  slots: (0 slots) slave
  replicates a84292141c294ad66fa0434c4e04c740769285e6
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001
  slots:0-6826,10923-12287 (8192 slots) master
  0 additional replica(s)
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003
  slots:12288-16383 (4096 slots) master
  2 additional replica(s)
S: 3b47db46737aeb0d070f490f404fbf9d6a6f4dae 10.1.5.205:7004
  slots: (0 slots) slave
  replicates faa7e2b785c07ed4b074b485cd78072e76337372
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006
  slots: (0 slots) slave
  replicates faa7e2b785c07ed4b074b485cd78072e76337372
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
[root@hadoop05 7004redis]#

```

### 三、从节点切换

新节点加入时，如果未指定，默认是主节点，且无 slot 分配。

```
[root@hadoop05 7004redis]#  
[root@hadoop05 7004redis]# redis-trib.rb check 10.1.5.205:7002  
>>> Performing Cluster Check (using node 10.1.5.205:7002)  
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002  
  slots:6827-10922 (4096 slots) master  
  1 additional replica(s)  
S: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005  
  slots: (0 slots) slave  
  replicates a84292141c294ad66fa0434c4e04c740769285e6  
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001  
  slots:0-6826,10923-12287 (8192 slots) master  
  0 additional replica(s)  
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003  
  slots:12288-16383 (4096 slots) master  
  1 additional replica(s)  
M: 3b47db46737aeb0d070f490f404fbf9d6a6f4dae 10.1.5.205:7004  
  slots: (0 slots) master  
  0 additional replica(s)  
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006  
  slots: (0 slots) slave  
  replicates faa7e2b785c07ed4b074b485cd78072e76337372  
[OK] All nodes agree about slots configuration.  
>>> check for open slots...  
>>> check slots coverage...  
[OK] All 16384 slots covered.  
[root@hadoop05 7004redis]#
```



此时可以通过 CLUSTER REPLICATE 命令进去指定为某个节点的从。

直连 7004,

**CLUSTER REPLICATE da150342ca272529d5dab62a557da5206305c333**

其中 da150342ca272529d5dab62a557da5206305c333 为 7001 的 id

```

[root@hadoop05 7004redis]#
[root@hadoop05 7004redis]# redis-cli -p 7004
127.0.0.1:7004>
127.0.0.1:7004>
127.0.0.1:7004> CLUSTER REPLICATE da150342ca272529d5dab62a557da5206305c333
OK
127.0.0.1:7004>
127.0.0.1:7004> cluster info
cluster_state:ok
cluster_slots_assigned:16384
cluster_slots_ok:16384
cluster_slots_pfail:0
cluster_slots_fail:0
cluster_known_nodes:6
cluster_size:3
cluster_current_epoch:16
cluster_my_epoch:16
cluster_stats_messages_sent:140
cluster_stats_messages_received:140
127.0.0.1:7004> quit
[root@hadoop05 7004redis]# redis-trib.rb check 10.1.5.205:7002
>>> Performing Cluster Check (using node 10.1.5.205:7002)
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002
  slots:6827-10922 (4096 slots) master
  1 additional replica(s)
S: cf285ac33cd79f2bfd604411185d15744ad5f4cf0 10.1.5.205:7005
  slots: (0 slots) slave
  replicates a84292141c294ad66fa0434c4e04c740769285e6
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001
  slots:0-6826,10923-12287 (8192 slots) master
  1 additional replica(s)
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003
  slots:12288-16383 (4096 slots) master
  1 additional replica(s)
S: 3b47db46737aeb0d070f490f404fbf9d6a6f4dae 10.1.5.205:7004
  slots: (0 slots) slave
  replicates da150342ca272529d5dab62a557da5206305c333
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006
  slots: (0 slots) slave
  replicates faa7e2b785c07ed4b074b485cd78072e76337372
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.

```

同样，即使 7004 已经有 master，也可以重新指定其他 master 代替当前 master，比如

下图，重新指定后,7003 有两个 slave

7001 master -->无 slave

7002 master --> 7005 slave

7003 master --> 7004/7006slave

注意：这种操作对初始创建集群时分配不合理，而又不能停止服务，需要重新调整时特别

有用。



```

[root@hadoop05 7004redis]#
[root@hadoop05 7004redis]#
[root@hadoop05 7004redis]# redis-cli -p 7004
127.0.0.1:7004>
127.0.0.1:7004>
127.0.0.1:7004>
127.0.0.1:7004> CLUSTER REPLICATE faa7e2b785c07ed4b074b485cd78072e76337372
OK
127.0.0.1:7004> quit
[root@hadoop05 7004redis]# redis-trib.rb check 10.1.5.205:7002
>>> Performing Cluster check (using node 10.1.5.205:7002)
M: a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002
slots:6827-10922 (4096 slots) master
1 additional replica(s)
S: cf285ac33cd79f2bfd604411185d15744ad5fcf0 10.1.5.205:7005
slots: (0 slots) slave
replicates a84292141c294ad66fa0434c4e04c740769285e6
M: da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001
slots:0-6826,10923-12287 (8192 slots) master
0 additional replica(s)
M: faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003
slots:12288-16383 (4096 slots) master
2 additional replica(s)
S: 3b47db46737aeb0d070f490f404fbf9d6a6f4dae 10.1.5.205:7004
slots: (0 slots) slave
replicates faa7e2b785c07ed4b074b485cd78072e76337372
S: 3d81d7eda6619c1deaa63fab6dfde9fe3ac71852 10.1.5.205:7006
slots: (0 slots) slave
replicates faa7e2b785c07ed4b074b485cd78072e76337372
[OK] All nodes agree about slots configuration.
>>> check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.

```

使用 info 也可以查看到主从相关信息。

```

[root@hadoop05 7004redis]# redis-trib.rb info 10.1.5.205:7002
10.1.5.205:7002 (a8429214...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7001 (da150342...) -> 0 keys | 8192 slots | 0 slaves.
10.1.5.205:7003 (faa7e2b7...) -> 0 keys | 4096 slots | 2 slaves.
[OK] 0 keys in 3 masters.
0.00 keys per slot on average.

```



## 四 观察主从切换

停止前:

```
127.0.0.1:7001> cluster nodes
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531131093146 6 connected
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531131094149 4 connected
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531131095151 3 connected 10923-16383
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 1 connected 0-5460
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531131092144 5 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531131094148 2 connected 5461-10922
127.0.0.1:7001>
```

kill 掉 7001 进程:

```
127.0.0.1:7002>
127.0.0.1:7002> cluster nodes
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531131282803 6 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 myself,master - 0 0 2 connected 5461-10922
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531131286839 3 connected 10923-16383
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 master - 0 1531131285816 7 connected 0-5460
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531131284813 5 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 master,fail - 1531131261836 1531131258610 1 disconnected
127.0.0.1:7002>
```

发现 7001 处于 fail 状态, 且 7001 的从 7004 变成了主。使用 redis-trib.rb 查看也提示

7004 有 0 个额外的副本, 且另两个 master 都有 1 个额外的副本。

```
[root@hadoop05 7003redis]#
[root@hadoop05 7003redis]# redis-trib.rb check 127.0.0.1:7006
>>> Performing Cluster check (using node 127.0.0.1:7006)
M: a3bc24765d7c8e1d5de86db2978c091fd81758c9 127.0.0.1:7006
  slots:10923-16383 (5461 slots) master
  1 additional replica(s)
S: 72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005
  slots: (0 slots) slave
  replicates 5289506a768c2253c10c91311b7135e5bee86bbe
S: 10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003
  slots: (0 slots) slave
  replicates a3bc24765d7c8e1d5de86db2978c091fd81758c9
M: a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004
  slots:0-5460 (5461 slots) master
  0 additional replica(s)
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
  slots:5461-10922 (5462 slots) master
  1 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
[root@hadoop05 7003redis]#
```

重启 7001

```
127.0.0.1:7002> cluster nodes
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531131398406 6 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 myself,master - 0 0 2 connected 5461-10922
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531131401413 3 connected 10923-16383
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 master - 0 1531131402415 7 connected 0-5460
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531131399409 5 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 slave a627628c4bd24c1e34454ff67cafe86858ccc84f 0 1531131400411 7 connected
127.0.0.1:7002>
```

发现 7004 仍然是主, 但 7001 变成了 7004 的从

备注:

a:如果集群任意 master 挂掉,且当前 master 没有 slave,集群进入 fail 状态,也可以理解成

进群 slot 映射[0-16383]不完全时，集群进入 fail 状态，当有 slave 进入时，自动转换为 master。

b:如果进群超过半数以上 master 挂掉，无论是否有 slave 集群都将进入 fail 状态。所以，如果三个 master 节点中，两个 master 挂掉，集群不可恢复，无论有多少个 slave。

c:当对应的 master、slave 同时挂时，集群不可通过新增节点恢复，除非修改所有点的 nodes.conf 信息，这种现象概率很低，当然也需要管理员监控好集群状态，防止这种情况出现。

## 五 观察主从切换时，数据是否丢失

对于 key:cn, value:abc。开始存储在 7001 上，当 7001 停止后，发现 7004 成了

master，且该 key 可以在 7004 上查到。

```
[root@hadoop05 redis-cluster]#  
[root@hadoop05 redis-cluster]# ./redis-3.2.3/src/redis-cli -h 10.1.5.205 -p 7005 -c  
10.1.5.205:7005> get cn  
-> Redirected to slot [4881] located at 10.1.5.205:7001  
"abc"  
10.1.5.205:7001> get cn  
Could not connect to Redis at 10.1.5.205:7001: connection refused  
not connected> quit  
[root@hadoop05 redis-cluster]# ./redis-3.2.3/src/redis-cli -h 10.1.5.205 -p 7005 -c  
10.1.5.205:7005> get cn  
-> Redirected to slot [4881] located at 10.1.5.205:7004  
"abc"  
10.1.5.205:7004>  
10.1.5.205:7004>  
10.1.5.205:7004> quit  
[root@hadoop05 redis-cluster]#  
[root@hadoop05 redis-cluster]#  
[root@hadoop05 redis-cluster]# ./redis-3.2.3/src/redis-cli -h 10.1.5.205 -p 7005 -c  
10.1.5.205:7005>  
10.1.5.205:7005>  
10.1.5.205:7005> cluster nodes  
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 master - 0 1531829993763 12 connected 10923-16383  
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 myself,slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 0 5 connected  
10518d23d2a78052c1c164ee68c197797631bf28 10.1.5.205:7003 slave a3bc24765d7c8e1d5de86db2978c091fd81758c9 0 1531829992759 12 connecte  
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531829994767 2 connected 5461-10922  
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 master - 0 1531829989744 13 connected 0-5460  
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 master,fail - 1531829862713 1531829857906 11 disconnected  
10.1.5.205:7005>
```

## 六、集群中添加和删除节点

### 6.1 cluster 方法

使用 redis-cli，连接，-c 与否则均可。

**cluster meet <ip> <port>**：将 ip 和 port 所指定的节点添加到集群当中，让它成为集群的一份子，nodes.conf 文件也会变化。

**cluster forget <node\_id>**：从集群中移除 node\_id 指定的节点，nodes.conf 文件中不会删除。

```
127.0.0.1:7001> cluster nodes
a627628c4bd24c1e34454ff67cfe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531898281167 22 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898284176 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531898283172 13 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898280165 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531898283674 21 connected
127.0.0.1:7001>
127.0.0.1:7001>
127.0.0.1:7001> cluster meet 10.1.5.205 7007
OK
127.0.0.1:7001>
127.0.0.1:7001> cluster nodes
a627628c4bd24c1e34454ff67cfe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531898299209 22 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898300211 21 connected 10923-16383
e83f6a5da7971ae1bcab378ae9f3f5ceb24ec18 10.1.5.205:7007 master - 0 1531898299310 0 connected
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531898298206 13 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898298708 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531898301213 21 connected
127.0.0.1:7001>
127.0.0.1:7001> cluster forget e83f6a5da7971ae1bcab378ae9f3f5ceb24ec18
OK
127.0.0.1:7001>
127.0.0.1:7001> cluster nodes
a627628c4bd24c1e34454ff67cfe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531898377423 22 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898379429 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531898378426 13 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898376420 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531898375417 21 connected
127.0.0.1:7001>
127.0.0.1:7001>
127.0.0.1:7001> cluster nodes
a627628c4bd24c1e34454ff67cfe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531898455618 22 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898456620 21 connected 10923-16383
e83f6a5da7971ae1bcab378ae9f3f5ceb24ec18 10.1.5.205:7007 master - 0 1531898453613 0 connected
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531898451610 13 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898452111 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531898454615 21 connected
127.0.0.1:7001>
```

注意：这种移除，其实并没有真正移除，因为各节点的 nodes.conf 中并没有删除该节点信息，只要该节点在线，稍等一会就会自动加入。

### 6.2 redis-trib.rb 方法

使用 redis-trib.rb 脚本

语法：

**add-node** new\_host:new\_port existing\_host:existing\_port

--slave，加入时默认是 master 节点。一个主可以有多个从。

--master-id <arg>

例如：redis-trib.rb add-node 10.1.5.205:7007 10.1.5.205:7002，

默认是 master 节点，且无 slot 分配，该操作也会直接修改 nodes.conf 文件。

```
[root@hadoop05 7001redis]# cat nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531898685239 22
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898689250 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531898684237 13
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898687245 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531898688248 21
vars currentEpoch 22 lastVoteEpoch 12
[root@hadoop05 7001redis]# redis-trib.rb add-node 10.1.5.205:7007 10.1.5.205:7002
>>> Adding node 10.1.5.205:7007 to cluster 10.1.5.205:7002
>>> Performing cluster check (using node 10.1.5.205:7002)
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
slots:5461-10922 (5462 slots) master
1 additional replica(s)
S: a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006
slots: (0 slots) slave
replicates 10518d23d2a78052c1c164ee68c197797651bf28
M: 10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003
slots:10923-16383 (5461 slots) master
1 additional replica(s)
S: a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
S: 72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005
slots: (0 slots) slave
replicates 5289506a768c2253c10c91311b7135e5bee86bbe
M: cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001
slots:0-5460 (5461 slots) master
1 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
>>> Send CLUSTER MEET to node 10.1.5.205:7007 to make it join the cluster.
[OK] New node added correctly.
[root@hadoop05 7001redis]# cat nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531898899030 22
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898902038 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531898903041 13
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898900034 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531898901036 21
fc4c63986c1f255a867382a8912bea49fdb5e68 10.1.5.205:7007 master - 0 1531898903144 0 connected
vars currentEpoch 22 lastVoteEpoch 12
[root@hadoop05 7001redis]#
```

添加节点时，使用--slave 选项，优先选择 slave 不足的主作为主

```
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531901343742 22
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531901341737 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531901346749 13
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531901345747 2 connected 5461-10922
vars currentEpoch 22 lastVoteEpoch 12
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]# redis-trib.rb add-node 10.1.5.205:7007 10.1.5.205:7002 --slave
[ERR] Wrong number of arguments for specified sub command
[root@hadoop05 redis-cluster]# redis-trib.rb add-node --slave 10.1.5.205:7007 10.1.5.205:7002
>>> Adding node 10.1.5.205:7007 to cluster 10.1.5.205:7002
>>> Performing cluster check (using node 10.1.5.205:7002)
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
slots:5461-10922 (5462 slots) master
1 additional replica(s)
M: 10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003
slots:10923-16383 (5461 slots) master
0 additional replica(s)
S: a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
S: 72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005
slots: (0 slots) slave
replicates 5289506a768c2253c10c91311b7135e5bee86bbe
M: cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001
slots:0-5460 (5461 slots) master
1 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
Automatically selected master 10.1.5.205:7003
>>> Send CLUSTER MEET to node 10.1.5.205:7007 to make it join the cluster.
waiting for the cluster to join.
>>> configure node as replica of 10.1.5.205:7003.
[OK] New node added correctly.
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531901455075 22
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531901457078 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531901456078 13
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531901458081 2 connected 5461-10922
908c580f3677939c299115ba3a9cbeebfb3c0ba 10.1.5.205:7007 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531901456881 21
vars currentEpoch 22 lastVoteEpoch 12
```

一个主可以有多个从，当从已经够时，自主协商哪一个作为主：



```
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531901619004 22 connected 0-5460
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531901457078 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531901615493 13 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531901458081 2 connected 5461-10922
908c580f3677939c299115ba3a9cbecbf3c0ba 10.1.5.205:7007 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531901616497 21 connected
vars currentEpoch 22 lastVoteEpoch 12
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]# redis-trib.rb add-node 10.1.5.205:7006 10.1.5.205:7002 --slave
[ERR] wrong number of arguments for specified sub command
[root@hadoop05 redis-cluster]# redis-trib.rb add-node --slave 10.1.5.205:7006 10.1.5.205:7002
>>> Adding node 10.1.5.205:7006 to cluster 10.1.5.205:7002
>>> Performing Cluster Check (using node 10.1.5.205:7002)
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
slots:5461-10922 (5462 slots) master
1 additional replica(s)
M: 10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003
slots:10923-16383 (5461 slots) master
1 additional replica(s)
S: a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
S: 72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005
slots: (0 slots) slave
replicates 5289506a768c2253c10c91311b7135e5bee86bbe
S: 908c580f3677939c299115ba3a9cbecbf3c0ba 10.1.5.205:7007
slots: (0 slots) slave
replicates 10518d23d2a78052c1c164ee68c197797651bf28
M: cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001
slots:0-5460 (5461 slots) master
1 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
Automatically selected master 10.1.5.205:7002
>>> Send CLUSTER MEET to node 10.1.5.205:7006 to make it join the cluster.
waiting for the cluster to join.
>>> Configure node as replica of 10.1.5.205:7002.
[OK] New node added correctly.
[root@hadoop05 redis-cluster]#
```

```
127.0.0.1:7001>
127.0.0.1:7001> cluster nodes
e89d78469c6800105c774e52f938f548cbe 10.1.5.205:7006 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531901619506 2 connected
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531901619004 22 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531901617498 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531901615493 13 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531901618503 2 connected 5461-10922
908c580f3677939c299115ba3a9cbecbf3c0ba 10.1.5.205:7007 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531901616497 21 connected
127.0.0.1:7001>
127.0.0.1:7001>
```

使用--master-id 选项，加入的为从节点时，指定主节点

```
[root@hadoop05 7006redis]#
[root@hadoop05 7006redis]# redis-trib.rb add-node --slave --master-id cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7006 10.1.5.205:7002
>>> Adding node 10.1.5.205:7006 to cluster 10.1.5.205:7002
>>> Performing Cluster Check (using node 10.1.5.205:7002)
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
slots:5461-10922 (5462 slots) master
1 additional replica(s)
M: 10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003
slots:10923-16383 (5461 slots) master
1 additional replica(s)
S: a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
S: 72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005
slots: (0 slots) slave
replicates 5289506a768c2253c10c91311b7135e5bee86bbe
S: 908c580f3677939c299115ba3a9cbecbf3c0ba 10.1.5.205:7007
slots: (0 slots) slave
replicates 10518d23d2a78052c1c164ee68c197797651bf28
M: cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001
slots:0-5460 (5461 slots) master
1 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
>>> Send CLUSTER MEET to node 10.1.5.205:7006 to make it join the cluster.
waiting for the cluster to join.
>>> Configure node as replica of 10.1.5.205:7001.
[OK] New node added correctly.
[root@hadoop05 7006redis]#
[root@hadoop05 7006redis]#
[root@hadoop05 7006redis]# cat nodes.conf
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 master - 0 1531902339260 22 connected 0-5460
ada7ed63edcc049a4c1ec30dbae58eb53c308b7 10.1.5.205:7006 myself,slave cef32f21f38fe366352ec43f3e271a221d34d712 0 0 0 connected
908c580f3677939c299115ba3a9cbecbf3c0ba 10.1.5.205:7007 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531902339361 21 connected
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531902339259 2 connected
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531902339259 21 connected 10923-16383
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531902339461 22 connected
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531902339159 2 connected 5461-10922
vars currentEpoch 22 lastVoteEpoch 0
[root@hadoop05 7006redis]#
```

删除节点：

语法：

del-node host:port node\_id

举例：

redis-trib.rb del-node 10.1.5.205:7007

fc4c63986c1f2355a867382a8912bea49fdb5e68

该操作会直接修改 nodes.conf，永久生效，和方法一不同

```
[root@hadoop05 7001redis]# cat nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 15318
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531898902038 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 15318
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531898900034 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 15318
fc4c63986c1f2355a867382a8912bea49fdb5e68 10.1.5.205:7007 master - 0 1531898903144 0 connected
vars currentEpoch 22 lastVoteEpoch 12
[root@hadoop05 7001redis]#
[root@hadoop05 7001redis]# redis-trib.rb del-node 10.1.5.205:7007 fc4c63986c1f2355a867382a8912bea49fdb5e68
>>> Removing node fc4c63986c1f2355a867382a8912bea49fdb5e68 from cluster 10.1.5.205:7007
>>> Sending CLUSTER FORGET messages to the cluster...
>>> SHUTDOWN the node.
[root@hadoop05 7001redis]# cat nodes.conf
a627628c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 15318
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 0-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531899584865 21 connected 10923-16383
72b17a468476e99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 15318
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531899580852 2 connected 5461-10922
a3bc24765d7c8e1d5de86db2978c091fd81758c9 10.1.5.205:7006 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 15318
vars currentEpoch 22 lastVoteEpoch 12
[root@hadoop05 7001redis]#
```



## 七、在线迁移 slot

```
reshard      host:port
             --from <arg>
             --to <arg>
             --slots <arg>
             --yes
             --timeout <arg>
             --pipeline <arg>
```

**host:port**: 这个是必传参数，用来从一个节点获取整个集群信息，相当于获取集群信息的入口。

**--from <arg>**: 需要从哪些源节点上迁移 slot，可从多个源节点完成迁移，以逗号隔开，传递的是节点的 node id，还可以直接传递--from all，这样源节点就是集群的所有节点，不传递该参数的话，则会在迁移过程中提示用户输入。

**--to <arg>**: slot 需要迁移的目的节点的 node id，目的节点只能填写一个，不传递该参数的话，则会在迁移过程中提示用户输入。

**--slots <arg>**: 需要迁移的 slot 数量，不传递该参数的话，则会在迁移过程中提示用户输入。

**--yes**: 设置该参数，可以在打印执行 reshard 计划的时候，提示用户输入 yes 确认后再执行 reshard。

**--timeout <arg>**: 设置 migrate 命令的超时时间。

**--pipeline <arg>**: 定义 cluster getkeysinslot 命令一次取出的 key 数量，不传的话使用默认值为 10。

比如：从 7001 移动 100 个槽位到 7002，默认是从小到大，100 就是 0-99，可使用--slots 指定移动的个数。

**举例：**

redis-trib.rb reshard --from cef32f21f38fe366352ec43f3e271a221d34d712 --  
to 5289506a768c2253c10c91311b7135e5bee86bbe --yes 10.1.5.205:7003

```
[root@hadoop05 ~]# redis-trib.rb reshard --from cef32f21f38fe366352ec43f3e271a221d34d712 --to 5289506a768c2253c10c91311b7135e5bee86bbe --yes 10.1.5.205:7003
>>> Performing cluster check (using node 10.1.5.205:7002)
M: 10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003
slots:10923-16383 (5461 slots) master
1 additional replica(s)
S: 72b17a46847ee99060861cbc8de60ff5aacbe50e 10.1.5.205:7005
slots: (0 slots) slave
replicates 5289506a768c2253c10c91311b7135e5bee86bbe
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
slots:5461-10922 (5462 slots) master
1 additional replica(s)
S: a62708c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
S: 908c580f36779339c299115ba3a9cbecbf3c0ba 10.1.5.205:7007
slots: (0 slots) slave
replicates 10518d23d2a78052c1c164ee68c197797651bf28
S: 69ff3efeb09d97f7372ea781fef52145d3b7ee 10.1.5.205:7006
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
M: cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001
slots:0-5460 (5461 slots) master
2 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
[OK] All 16384 slots covered.
How many slots do you want to move (from 1 to 16384)? 100
Ready to move 100 slots.
Source nodes:
M: cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001
slots:0-5460 (5461 slots) master
2 additional replica(s)
Destination node:
M: 5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002
slots:5461-10922 (5462 slots) master
```

如果不加--yes，系统会提示如下：

```
Moving slot 294 from cef32f21f38fe366352ec43f3e271a221d34d712 to 5289506a768c2253c10c91311b7135e5bee86bbe
Moving slot 295 from cef32f21f38fe366352ec43f3e271a221d34d712 to 5289506a768c2253c10c91311b7135e5bee86bbe
Moving slot 296 from cef32f21f38fe366352ec43f3e271a221d34d712 to 5289506a768c2253c10c91311b7135e5bee86bbe
Moving slot 297 from cef32f21f38fe366352ec43f3e271a221d34d712 to 5289506a768c2253c10c91311b7135e5bee86bbe
Moving slot 298 from cef32f21f38fe366352ec43f3e271a221d34d712 to 5289506a768c2253c10c91311b7135e5bee86bbe
Moving slot 299 from cef32f21f38fe366352ec43f3e271a221d34d712 to 5289506a768c2253c10c91311b7135e5bee86bbe
Do you want to proceed with the proposed reshard plan (yes/no)? yes
Moving slot 200 from 10.1.5.205:7001 to 10.1.5.205:7002:
Moving slot 201 from 10.1.5.205:7001 to 10.1.5.205:7002:
Moving slot 202 from 10.1.5.205:7001 to 10.1.5.205:7002:
```

移动后：

```
127.0.0.1:7001>
127.0.0.1:7001> cluster nodes
a62708c4bd24c1e34454ff67cafe86858ccc84f 10.1.5.205:7004 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531903920697 22 connected
cef32f21f38fe366352ec43f3e271a221d34d712 10.1.5.205:7001 myself,master - 0 0 22 connected 100-5460
10518d23d2a78052c1c164ee68c197797651bf28 10.1.5.205:7003 master - 0 1531903921699 21 connected 10923-16383
5289506a768c2253c10c91311b7135e5bee86bbe 10.1.5.205:7002 master - 0 1531903917187 23 connected 0-99 5461-10922
72b17a46847ee99060861cbc8de60ff5aacbe50e 10.1.5.205:7005 slave 5289506a768c2253c10c91311b7135e5bee86bbe 0 1531903918691 23 connected
908c580f36779339c299115ba3a9cbecbf3c0ba 10.1.5.205:7007 slave 10518d23d2a78052c1c164ee68c197797651bf28 0 1531903919693 21 connected
69ff3efeb09d97f7372ea781fef52145d3b7ee 10.1.5.205:7006 slave cef32f21f38fe366352ec43f3e271a221d34d712 0 1531903919192 22 connected
127.0.0.1:7001>
127.0.0.1:7001>
127.0.0.1:7001>
```

注意：to 的节点必须是 master，否则出现错误

```
slots: (0 slots) slave
replicates cef32f21f38fe366352ec43f3e271a221d34d712
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
How many slots do you want to move (from 1 to 16384)? 100
*** The specified node is not known or not a master, please retry.
[root@hadoop05 ~]#
```

## 八 平衡集群节点的 slot 数量

```
rebalance    host:port
             --weight <arg>
             --auto-weights
             --use-empty-masters
             --timeout <arg>
             --simulate
             --pipeline <arg>
             --threshold <arg>
```

host:port: 这个是必传参数，用来从一个节点获取整个集群信息，相当于获取集群信息的入口。

--weight <arg>: 节点的权重，格式为 node\_id=weight，如果需要为多个节点分配权重的话，需要添加多个--weight <arg>参数，即--weight b31e3a2e=5 --weight

60b8e3a1=5，node\_id 可为节点名称的前缀，只要保证前缀位数能唯一区分该节点即可。

没有传递--weight 的节点的权重默认为 1。

--auto-weights: 这个参数在 rebalance 流程中并未用到。

--threshold <arg>: 只有节点需要迁移的 slot 阈值超过 threshold，才会执行 rebalance 操作。具体计算方法可以参考下面的 rebalance 命令流程的第四步。

--use-empty-masters: rebalance 是否考虑没有节点的 master，默认没有分配 slot 节点的 master 是不参与 rebalance 的，设置--use-empty-masters 可以让没有分配 slot 的节点参与 rebalance，这个参数在扩容集群 master 数量时特别有用。

--timeout <arg>: 设置 migrate 命令的超时时间。

--simulate: 设置该参数，可以模拟 rebalance 操作，提示用户会迁移哪些 slots，而不会真正执行迁移操作。

--pipeline <arg>: 与 reshar 的 pipeline 参数一样，定义 cluster getkeysinslot 命令一

次取出的 key 数量，不传的话使用默认值为 10。

## 8.1 没有新 master 加入

平衡前：

```
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003 master - 0 1531969323639 8 connected 100-199 10923-16383
bc3e08e793344c02c665fe7200053f86f45bc9bd 10.1.5.205:7004 slave da150342ca272529d5dab62a557da5206305c333 0 1531969319605 4 connected
cf285ac33cd79f2bf6d04411185d15744ad5fcf0 10.1.5.205:7005 slave a84292141c294ad66fa0434c4e04c740769285e6 0 1531969324641 7 connected
da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001 myself,master - 0 0 1 connected 200-5460
a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002 master - 0 1531969322635 7 connected 0-99 5461-10922
3d81d7eda6619c1dea63fab6dfde9fe3ac71852 10.1.5.205:7006 slave faa7e2b785c07ed4b074b485cd78072e76337372 0 1531969321610 8 connected
vars currentEpoch 8 lastVoteEpoch 0
[root@hadoop05 redis-cluster]#
```

执行平衡命令：

redis-trib.rb rebalance 10.1.5.205:7001

```
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]# redis-trib.rb rebalance 10.1.5.205:7001
>>> Performing Cluster Check (using node 10.1.5.205:7001)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
>>> Rebalancing across 3 nodes. Total weight = 3
Moving 101 slots from 10.1.5.205:7002 to 10.1.5.205:7001
#####
Moving 100 slots from 10.1.5.205:7003 to 10.1.5.205:7001
#####
```

可以明显的看到从 7002 和 7003 各移动 100 个 slot 到 7001

平衡后：

```
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003 master - 0 1531969450955 8 connected 10923-16383
bc3e08e793344c02c665fe7200053f86f45bc9bd 10.1.5.205:7004 slave da150342ca272529d5dab62a557da5206305c333 0 1531969452963 9 connected
cf285ac33cd79f2bf6d04411185d15744ad5fcf0 10.1.5.205:7005 slave a84292141c294ad66fa0434c4e04c740769285e6 0 1531969451975 7 connected
da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001 myself,master - 0 0 9 connected 0-5461
a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002 master - 0 1531969449952 7 connected 5462-10922
3d81d7eda6619c1dea63fab6dfde9fe3ac71852 10.1.5.205:7006 slave faa7e2b785c07ed4b074b485cd78072e76337372 0 1531969448951 8 connected
vars currentEpoch 9 lastVoteEpoch 0
[root@hadoop05 redis-cluster]#
```

## 8.2 有新的 master 加入

多加一个 master 节点进行平衡：

平衡前：

处于不平衡状态，且新加入的 7007 是 master，且未分配任何 slot。

```
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
9987adaef51933673afe8f5d7c0ea292c94f2182 10.1.5.205:7007 master - 0 1531970538226 0 connected
faa7e2b785c07ed4b074b485cd78072e76337372 10.1.5.205:7003 master - 0 1531970539228 13 connected 0-299 10923-16383
bc3e08e793344c02c665fe7200053f86f45bc9bd 10.1.5.205:7004 slave da150342ca272529d5dab62a557da5206305c333 0 1531970541
cf285ac33cd79f2bf6d04411185d15744ad5fcf0 10.1.5.205:7005 slave a84292141c294ad66fa0434c4e04c740769285e6 0 1531970541
da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001 myself,master - 0 0 12 connected 600-5461
a84292141c294ad66fa0434c4e04c740769285e6 10.1.5.205:7002 master - 0 1531970542321 14 connected 300-599 5462-10922
3d81d7eda6619c1dea63fab6dfde9fe3ac71852 10.1.5.205:7006 slave faa7e2b785c07ed4b074b485cd78072e76337372 0 1531970540
vars currentEpoch 14 lastVoteEpoch 0
[root@hadoop05 redis-cluster]#
```

执行平衡命令：

redis-trib.rb rebalance --use-empty-masters 10.1.5.205:7001

可以看到从 7002 7003 7001 各移动了一定数量的 slot 到 7007

[illegible]

平衡后:

新增的 master 节点也接管了一部分 slot, 且 7007 无 slave

```
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]# cat 7001redis/nodes.conf
9987adaef51933673afe6f5d7c0ea29c2c94f2182 10.1.5.205:7007 master - 0 1531970596763 15 connected 0-1365 5462-6826 10923-12287
faa7e2b785c07ed4b074b485d78072e76337372 10.1.5.205:7003 master - 0 1531970598769 13 connected 12288-16383
bc3e08e793544c02c665fe71100553f86f45bc9bd 10.1.5.205:7004 slave da150342ca272529d5dab62a557da5206305c333 0 1531970599770 12 co
c285ac3cd3c79f2bf6d0441128d5f17544a5dcfcf0 10.1.5.205:7005 slave aa84292141c294ad66fa0434c4e04c7407692856e 0 1531970595459 14 co
da150342ca272529d5dab62a557da5206305c333 10.1.5.205:7001 myself,master - 0 12 12 connected 1366-5461
aa84292141c294ad66fa0434c4e04c7407692856e 10.1.5.205:7002 master - 0 1531970600770 14 connected 6827-10922
3d81d7eda6619c1dea63fab6dfde9f3ec3b71852 10.1.5.205:7006 slave faa7e2b785c07ed4b074b485cd78072e76337372 0 1531970597765 13 co
vars currentEpoch 15 lastVoteEpoch 0
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]#
[root@hadoop05 redis-cluster]# redis-trib.rb info 10.1.5.205:7001
10.1.5.205:7001 (da150342...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7007 (9987adae...) -> 0 keys | 4096 slots | 0 slaves.
10.1.5.205:7003 (faa7e2b7...) -> 0 keys | 4096 slots | 1 slaves.
10.1.5.205:7002 (aa842921...) -> 0 keys | 4096 slots | 1 slaves.
[OK] 0 keys in 4 masters,
0.00 keys per slot on average.
```